

1. Splunk Alert: High Count of ext_authz_error - Critical !!

1. Descripti on	Runs every 5mins, Takes last 5 mins of data. CLOSELY monitor the failures for next few minutes
Possible Reasons	• Call from TAG to TaapSec did not reach TaapSec • TaapSec took too long to process & respond • TaapSec responded within time but the response did not reach TAG within time • High Memory Utilization of the TAG Node
Alert Query	index=meg_dataplane respCodeDetails=ext_authz_error stats count by kubernetes_cluster search count > 500
Channel	API CFE - Alerts tag-prod-alert-action Microsoft Teams

Triaging steps:

1. **Respond right away to the Alert with** "Checking TAG & TaapSec Logs" **[No "Ack" please]**
 - a. Post below message in [API CFE - Internal](#) / [tag-internal](#) / [Microsoft Teams](#) stating "Received alert for ext_authz_error" - <Teams link of the alert> and use the tag @TAG T2 (PE) to notify T2 team of platform engineering.
 - b. If two consecutive alerts are received for this please engage the oncall (dpo core team member) and start a bridge to engage TAG & TaapSec Tier-2 team member as documented [here](#) (Row #3)
2. Execute below query for for 1min span in splunk
 - a. [index=meg_dataplane respCodeDetails=ext_authz_error | timechart span=1m count by kubernetes_cluster](#)
3. To get the necessary info for cluster, namespace and container you can use query
 - a. [index=meg_dataplane respCodeDetails=ext_authz_error | stats count by kubernetes_cluster,kubernetes_namespace,kubernetes_deployment_name](#)
4. Check if the cluster hosted on TKE, DUCK/Conduktor*. Accordingly check the Grafana dashboard (1) [Application Monitoring Dashboard](#) or [Resource Usage Splunk O11y Dashboard](#) (right now splunk Olly has only TKE clusters listed), filter with the values from the previous query:

Grafana Dashboard Filter	Splunk Field
--------------------------	--------------

Cluster	kubernetes_cluster
namespace	kubernetes_namespace
Deployment/Statefulset/Daemonset/Jobs	kubernetes_deployment_name

***Note:** To know if a cluster is TKE or Conduktor, as a general rule, any cluster name starting with px or tt is TKE and you can assume the rest are Conduktor, but you can refer [here](#) for a list of clusters and namespaces grouped into tedge (TKE) and aedge (Conduktor)


4. For the Memory & Usage of the Kubernetes Cluster and namespace in which errors are reported (based on step#2)

a. **If the % Value for Memory usage is > 65%, PLEASE CALL THE TAG PE Tier-2 Team**

b. **If the % Value for CPU usage is > 65%, PLEASE CALL THE TAG PE Tier-2 Team**

c. **If you can see pods are restarting, PLEASE CALL THE TAG PE Tier-2 Team**

d. Please update the thread with the screenshot of the top portion of the dashboard and the link to it in [API CFE - Internal](#) | [tag-internal](#) | [Microsoft Teams](#)



5. Execute Below Splunk queries to check if there are any issues with **TaapSec Service** and update the thread with the information from below queries, if no errors are found when running this, have an update mentioning there are no 5xx errors / connection errors on taap validator ms, otherwise update with the splunk link and a screenshot of the table

a. [index=meg_dataplane serviceName=taap-validator-ms "java.net.ConnectException" | stats count by kubernetes_cluster, kubernetes_pod_id, kubernetes_pod_name](#)

b. `index=meg_dataplane serviceName=taap-validator-ms httpStatusCode="5*" | stats count by kubernetes_cluster, kubernetes_pod_id, kubernetes_pod_name`

6. Check some transaction ids from the queries (based on step#2) and replace a couple of them in:

a. `index=meg_dataplane sourcetype=taapsec REPLACE_TRANSACTION_ID | table kubernetes_deployment_name`


7. Check if the cluster hosted on TKE, DUCK / Conduktor. Accordingly check the Grafana dashboard (1) [Application Monitoring Dashboard](#) or [Resource Usage Splunk O11y Dashboard](#) (right now splunk Olly has only TKE clusters listed), select the platform TKE or Conduktor, the cluster (from step#2), the namespace (from step#2) and the deployment (from step#5) and check if any spikes on Memory usage, CPU usage or Network errors / dropped packets.

a. **If the % Value for Memory usage is > 65%, PLEASE CALL THE TAAPSEC Tier-2 Team**

b. **If the % Value for CPU usage is > 65%, PLEASE CALL THE TAAPSEC Tier-2 Team**

c. **If you can see pods are restarting, PLEASE CALL THE TAAPSEC Tier-2 Team**

d. Please update the thread with the screenshot of the top portion of the dashboard and the link to it in [API CFE - Internal](#) | [tag-internal](#) | [Microsoft Teams](#)



8. For `ext_authz_error` alerts, validate whether failures are due to long-running client requests (~7 seconds) by checking for entries with `clientDuration > 6500ms` using:

`index=meg_dataplane respCodeDetails=ext_authz_error`

`kubernetes_cluster=REPLACE_CLUSTER_NAME`

`kubernetes_namespace=REPLACE_NAMESPACE proxyName=* clientDuration>6500`

Then confirm whether these delays correlate with node-specific latency in Taap-validator by checking for `processingTimeMillis > 1000ms` and visualizing average node performance:

```
index=meg_dataplane eventType=performAuthChecks
kubernetes_namespace=REPLACE_NAMESPACE kubernetes_cluster IN
(REPLACE_CLUSTER_NAME) kubernetes_deployment_name=taap-validator-ms-prd*
processingTimeMillis>1000 | timechart span=1s avg(processingTimeMillis) by
kubernetes_node_name
```

If high latency appears only on specific nodes and returns to normal shortly afterward, this typically indicates a node migration event (e.g., VMotion) and not a system-wide issue.

9. **If it is just a single spike which doesn't last for more than 5 mins and there are no more errors**, then no need to call the Tier-2 team
10. Please raise a P4 Incident and assign it to the relevant teams for tracking the RCA. Please refer [📖 27. INC Assignment Group](#) for more details.